

Exploratory Data Analysis (EDA) & Data Cleaning for House Pricing Dataset

Data Loading & Initial Analysis

```
In [1]: # Importing necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: # Load the dataset
```

```
In [3]: file_path = "raw_house_data.csv"
df = pd.read_csv(file_path)
```

```
In [4]: # Initial exploration
```

```
In [5]: print("Initial Dataset Info:")
print(df.info())
```

```
Initial Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MLS                    5000 non-null   int64
1   sold_price             5000 non-null   float64
2   zipcode                5000 non-null   int64
3   longitude              5000 non-null   float64
4   latitude               5000 non-null   float64
5   lot_acres              4990 non-null   float64
6   taxes                  5000 non-null   float64
7   year_built             5000 non-null   int64
8   bedrooms               5000 non-null   int64
9   bathrooms              4994 non-null   float64
10  sqrt_ft                4944 non-null   float64
11  garage                 4993 non-null   float64
12  kitchen_features       4967 non-null   object
13  fireplaces             5000 non-null   object
14  floor_covering         4999 non-null   object
15  HOA                    4438 non-null   object
dtypes: float64(8), int64(4), object(4)
memory usage: 625.1+ KB
None
```

```
In [6]: df.head()
```

Out[6]:

	MLS	sold_price	zipcode	longitude	latitude	lot_acres	taxes	year_built
--	-----	------------	---------	-----------	----------	-----------	-------	------------

0	21530491	5300000.0	85637	-110.378200	31.356362	2154.00	5272.00	1941
---	----------	-----------	-------	-------------	-----------	---------	---------	------

1	21529082	4200000.0	85646	-111.045371	31.594213	1707.00	10422.36	1997
---	----------	-----------	-------	-------------	-----------	---------	----------	------

2	3054672	4200000.0	85646	-111.040707	31.594844	1707.00	10482.00	1997
---	---------	-----------	-------	-------------	-----------	---------	----------	------

3	21919321	4500000.0	85646	-111.035925	31.645878	636.67	8418.58	1930
---	----------	-----------	-------	-------------	-----------	--------	---------	------

4	21306357	3411450.0	85750	-110.813768	32.285162	3.21	15393.00	1995
---	----------	-----------	-------	-------------	-----------	------	----------	------



In [7]: `df.tail()`

Out[7]:

	MLS	sold_price	zipcode	longitude	latitude	lot_acres	taxes	year_built
--	-----	------------	---------	-----------	----------	-----------	-------	------------

4995	21810382	495000.0	85641	-110.661829	31.907917	4.98	2017.00	2000
------	----------	----------	-------	-------------	-----------	------	---------	------

4996	21908591	550000.0	85750	-110.858556	32.316373	1.42	4822.01	1990
------	----------	----------	-------	-------------	-----------	------	---------	------

4997	21832452	475000.0	85192	-110.755428	32.964708	12.06	1000.00	1990
------	----------	----------	-------	-------------	-----------	-------	---------	------

4998	21900515	550000.0	85745	-111.055528	32.296871	1.01	5822.93	2000
------	----------	----------	-------	-------------	-----------	------	---------	------

4999	4111490	450000.0	85621	-110.913054	31.385259	4.16	2814.48	1990
------	---------	----------	-------	-------------	-----------	------	---------	------



In [8]: `df.shape`

```
Out[8]: (5000, 16)
```

```
In [9]: df.count()
```

```
Out[9]: MLS                5000
sold_price                5000
zipcode                  5000
longitude                 5000
latitude                  5000
lot_acres                 4990
taxes                     5000
year_built                5000
bedrooms                  5000
bathrooms                 4994
sqrt_ft                   4944
garage                    4993
kitchen_features          4967
fireplaces                5000
floor_covering             4999
HOA                       4438
dtype: int64
```

```
In [10]: df.describe()
```

```
Out[10]:
```

	MLS	sold_price	zipcode	longitude	latitude	lot_acres
count	5.000000e+03	5.000000e+03	5000.000000	5000.000000	5000.000000	4990.000000
mean	2.127070e+07	7.746262e+05	85723.025600	-110.912107	32.308512	4.661311
std	2.398508e+06	3.185556e+05	38.061712	0.120629	0.178028	51.685230
min	3.042851e+06	1.690000e+05	85118.000000	-112.520168	31.356362	0.000000
25%	2.140718e+07	5.850000e+05	85718.000000	-110.979260	32.277484	0.580000
50%	2.161469e+07	6.750000e+05	85737.000000	-110.923420	32.318517	0.990000
75%	2.180480e+07	8.350000e+05	85749.000000	-110.859078	32.394334	1.757500
max	2.192856e+07	5.300000e+06	86323.000000	-109.454637	34.927884	2154.000000

Data Cleaning

Data Type Conversion

```
In [11]: convert_cols = ["bathrooms", "sqrt_ft", "garage", "fireplaces", "HOA"]
for col in convert_cols:
    df[col] = pd.to_numeric(df[col], errors='coerce')
```


Handle missing values

```
In [12]: df.fillna({"lot_acres":df["lot_acres"].median()}, inplace=True)
df.fillna({"sqrt_ft":df["sqrt_ft"].median()}, inplace=True)
```

```
df.fillna({"garage": 0}, inplace=True)
df.fillna({"HOA": 0}, inplace=True)
```


```
In [13]: sns.boxplot(x=df['sold_price'])
```

```
Out[13]: <Axes: xlabel='sold_price'>
```



```
In [14]: sns.boxplot(x=df['HOA'])
```

```
Out[14]: <Axes: xlabel='HOA'>
```



Outlier Removal using IQR method

```
In [15]: column = df.select_dtypes(include=['float64', 'int64']).columns
```

```
In [16]: def remove_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]
```

```
In [17]: outlier_cols = ["lot_acres", "sold_price", "taxes", "sqrt_ft", "HOA"]
    for col in outlier_cols:
        df = remove_outliers(df, col)
```

```
In [18]: # Summary Statistics After Cleaning:
```

```
In [19]: print("\nSummary Statistics After Cleaning:")
    print(df.describe())
```

Summary Statistics After Cleaning:

	MLS	sold_price	zipcode	longitude	latitude \
count	3.733000e+03	3.733000e+03	3733.000000	3733.000000	3733.000000
mean	2.135377e+07	6.837639e+05	85726.135548	-110.916177	32.324772
std	2.101180e+06	1.362301e+05	32.688514	0.092688	0.134497
min	3.042851e+06	3.750000e+05	85118.000000	-111.430863	31.458609
25%	2.140831e+07	5.750000e+05	85718.000000	-110.975535	32.285978
50%	2.161784e+07	6.500000e+05	85737.000000	-110.922752	32.319066
75%	2.180678e+07	7.500000e+05	85750.000000	-110.861144	32.396889
max	2.192856e+07	1.185000e+06	85935.000000	-109.861617	34.314889

	lot_acres	taxes	year_built	bedrooms	bathrooms \
count	3733.000000	3733.000000	3733.000000	3733.000000	3729.000000
mean	1.06910	6006.996552	1992.924993	3.850522	3.600697
std	0.84015	2092.235652	49.317624	0.826379	0.968307
min	0.00000	459.530000	0.000000	2.000000	2.000000
25%	0.51000	4708.000000	1987.000000	3.000000	3.000000
50%	0.87000	5945.000000	1999.000000	4.000000	4.000000
75%	1.20000	7272.620000	2005.000000	4.000000	4.000000
max	3.50000	11809.000000	2019.000000	18.000000	36.000000

	sqrt_ft	garage	fireplaces	HOA
count	3733.000000	3733.000000	3714.000000	3733.000000
mean	3423.058666	2.737209	1.721055	63.131372
std	617.383578	0.918898	1.000271	66.155298
min	1780.000000	0.000000	0.000000	0.000000
25%	2998.000000	2.000000	1.000000	3.000000
50%	3401.000000	3.000000	2.000000	44.000000
75%	3811.000000	3.000000	2.000000	100.000000
max	5125.000000	12.000000	8.000000	263.000000

In [20]: df.head()

Out[20]: **LS** **sold_price** **zipcode** **longitude** **latitude** **lot_acres** **taxes** **year_built** **bedroom**

40 1125000.0 85718 -110.883547 32.329763 1.33 8654.00 1986

37 1100000.0 85750 -110.866891 32.321968 1.17 6565.93 1994

50 1180000.0 85750 -110.868487 32.316324 1.30 9590.16 1993

55 1175500.0 85718 -110.940650 32.347873 1.23 11674.00 2004

03 1125478.0 85755 -110.973498 32.460529 1.71 3171.39 2017

```
In [21]: #df = df.drop(columns=['kitchen_features', 'floor_covering'])
```

```
In [22]: print(df)
```

	MLS	sold_price	zipcode	longitude	latitude	lot_acres	\
398	21329440	1125000.0	85718	-110.883547	32.329763	1.33	
400	21500337	1100000.0	85750	-110.866891	32.321968	1.17	
411	21206450	1180000.0	85750	-110.868487	32.316324	1.30	
412	21224755	1175500.0	85718	-110.940650	32.347873	1.23	
428	21703603	1125478.0	85755	-110.973498	32.460529	1.71	
...	
4992	3056450	525000.0	85614	-110.980945	31.824287	3.01	
4993	21908358	565000.0	85750	-110.820216	32.307646	0.83	
4994	21909379	535000.0	85718	-110.922291	32.317496	0.18	
4996	21908591	550000.0	85750	-110.858556	32.316373	1.42	
4998	21900515	550000.0	85745	-111.055528	32.296871	1.01	

	taxes	year_built	bedrooms	bathrooms	sqrt_ft	garage	\
398	8654.00	1986	4	5.0	5023.0	3.0	
400	6565.93	1994	4	4.0	3870.0	3.0	
411	9590.16	1993	4	3.0	5029.0	3.0	
412	11674.00	2004	4	5.0	4143.0	3.0	
428	3171.39	2017	3	4.0	3436.0	3.0	
...	
4992	5122.84	2007	3	3.0	3512.0	3.0	
4993	4568.71	1986	4	3.0	2813.0	2.0	
4994	4414.00	2002	3	2.0	2106.0	2.0	
4996	4822.01	1990	4	3.0	2318.0	3.0	
4998	5822.93	2009	4	4.0	3724.0	3.0	

	kitchen_features	fireplaces	\
398	Compactor, Dishwasher, Garbage Disposal, Refri...	3.0	
400	Dishwasher, Garbage Disposal, Refrigerator, Mi...	2.0	
411	Dishwasher, Garbage Disposal, Refrigerator, Mi...	3.0	
412	Dishwasher, Garbage Disposal, Refrigerator, Mi...	1.0	
428	Dishwasher, Garbage Disposal	1.0	
...	
4992	Dishwasher, Garbage Disposal, Gas Range, Refri...	1.0	
4993	Dishwasher, Double Sink, Electric Range, Garba...	2.0	
4994	Dishwasher, Double Sink, Electric Range, Garba...	1.0	
4996	Dishwasher, Double Sink, Electric Range, Garba...	1.0	
4998	Dishwasher, Double Sink, Garbage Disposal, Gas...	1.0	

	floor_covering	HOA
398	Carpet, Natural Stone, Wood	179.0
400	Carpet, Natural Stone	58.0
411	Carpet, Ceramic Tile	40.0
412	Carpet, Ceramic Tile	159.0
428	Carpet, Natural Stone	0.0
...
4992	Concrete, Other: Cork	37.0
4993	Carpet, Mexican Tile	6.0
4994	Ceramic Tile	198.0
4996	Carpet, Ceramic Tile	43.0
4998	Carpet, Ceramic Tile	0.0

[3733 rows x 16 columns]

Save cleaned dataset

```
In [23]: df.to_csv("cleaned_house_data.csv", index=False)
print("\nCleaned dataset saved as 'cleaned_house_data.csv'.")
```

Cleaned dataset saved as 'cleaned_house_data.csv'.


Data Visualization

Bar Chart

Used to compare categorical values such as sold price, lot acres.

```
In [24]: df.sold_price.value_counts().nlargest(40).plot(kind='bar', figsize=(10,5))
```

Out[24]: <Axes: xlabel='sold_price'>



```
df.lot_acres.value_counts().nlargest(40).plot(kind='bar', figsize=(10,5))
```


Histogram of Sold Prices

- House prices were right-skewed, meaning some high-value properties affected the average.
- Suggested log transformation to normalize data for better model accuracy.

```
In [25]: plt.figure(figsize=(10, 5))
sns.histplot(df["sold_price"], bins=50, kde=True) #Kernel Density Estimation
plt.title("Distribution of Sold Prices")
plt.xlabel("Sold Price")
```



```
plt.ylabel("Count")
plt.show()
```

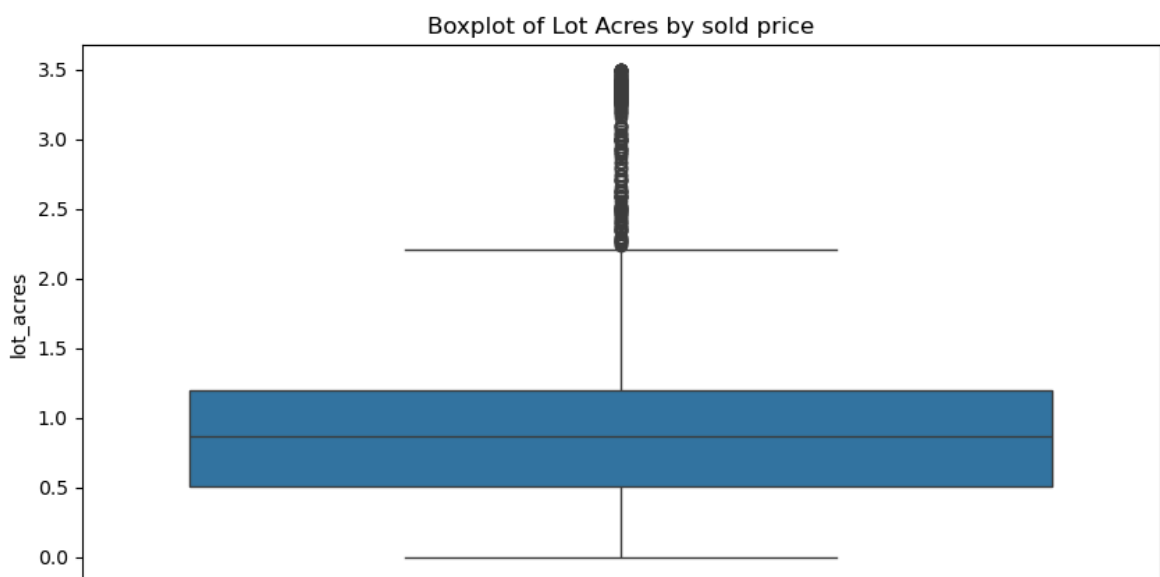


Boxplot of Lot Acres

Shows data distribution & outliers using quartiles.

- Before cleaning, several extreme outliers were identified in house prices.
- After applying IQR-based filtering, the data became more consistent and normally distributed.

```
In [26]: plt.figure(figsize=(10, 5))
sns.boxplot(y=df["lot_acres"])
plt.title("Boxplot of Lot Acres by sold price")
plt.show()
```




Correlation Heatmap

Helps identify feature relationships to select relevant predictors for modeling.

- Square footage, number of rooms, and overall quality had the highest correlation with house prices.
- Lot size had weak correlation, indicating it might not be a strong predictor.


```
In [27]: plt.figure(figsize=(12, 8))
sns.heatmap(df[column].corr(), annot=True, cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()
```



Scatter plot of Square Footage vs Price

A scatter plot is used to visualize the relationship between house size (square footage) and sold price. This helps identify whether larger homes tend to have higher prices.


```
In [28]: plt.figure(figsize=(10, 5))
sns.scatterplot(x=df["sqrt_ft"], y=df["sold_price"])
plt.title("Square Footage vs Sold Price")
plt.xlabel("Square Footage")
plt.ylabel("Sold Price")
plt.show()
```



Count Plot of Number of Bedrooms

A count plot is used to visualize the distribution of houses based on the number of bedrooms. This helps identify the most common house types in the dataset.

```
In [29]: plt.figure(figsize=(10, 5))
sns.countplot(x=df["bedrooms"])
plt.title("Count of Houses by Number of Bedrooms")
plt.xlabel("Bedrooms")
plt.ylabel("Count")
plt.show()
```

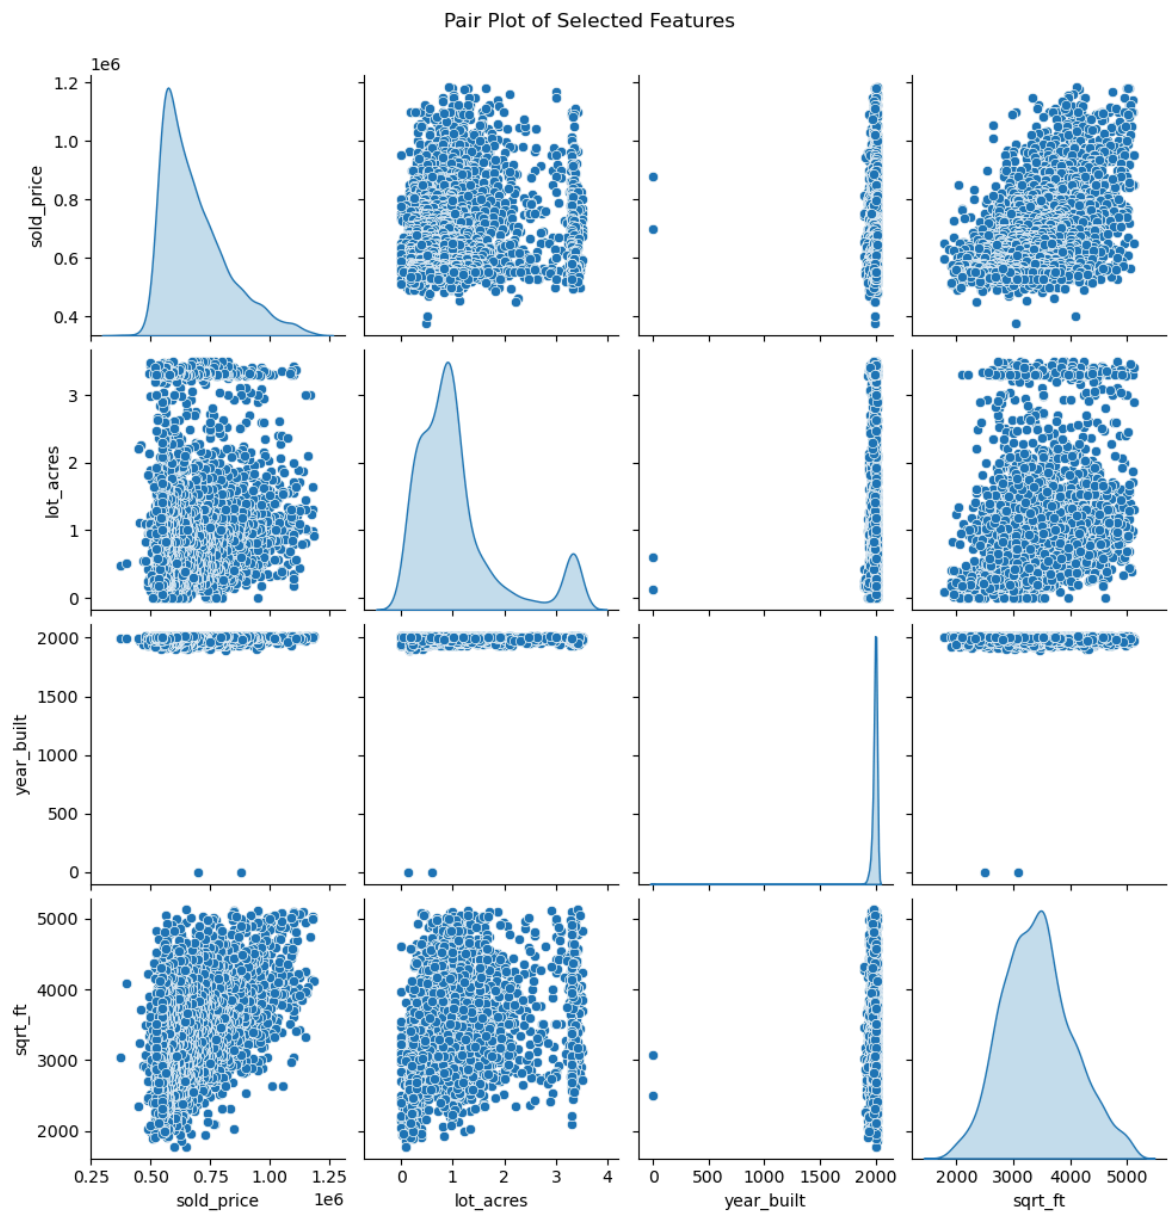


PairPlot of Selected Features

Displays pairwise relationships between multiple numerical variables.

- House price vs. square footage showed a strong positive correlation (larger homes tend to be more expensive).
- Year built vs. price suggested that newer homes generally have higher prices.
- Outliers were visible in some features like lot size.

```
In [30]: selectedColumn = ['sold_price', 'lot_acres', 'year_built', 'sqrt_ft']
sns.pairplot(df[selectedColumn], diag_kind="kde")
plt.suptitle("Pair Plot of Selected Features", y=1.02)
plt.show()
```



```
In [ ]:
```